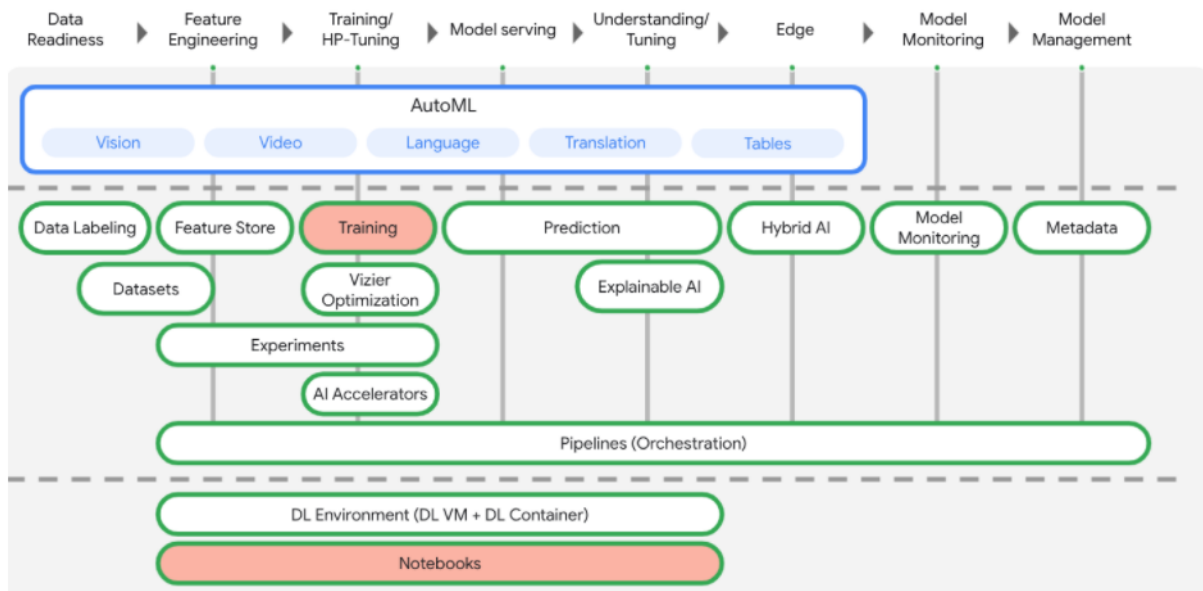


a) Vertex AI to run a hyperparameter tuning job for a TensorFlow model  
Reference: [https://codelabs.developers.google.com/vertex\\_hyperparameter\\_tuning#0](https://codelabs.developers.google.com/vertex_hyperparameter_tuning#0)

Objectives:

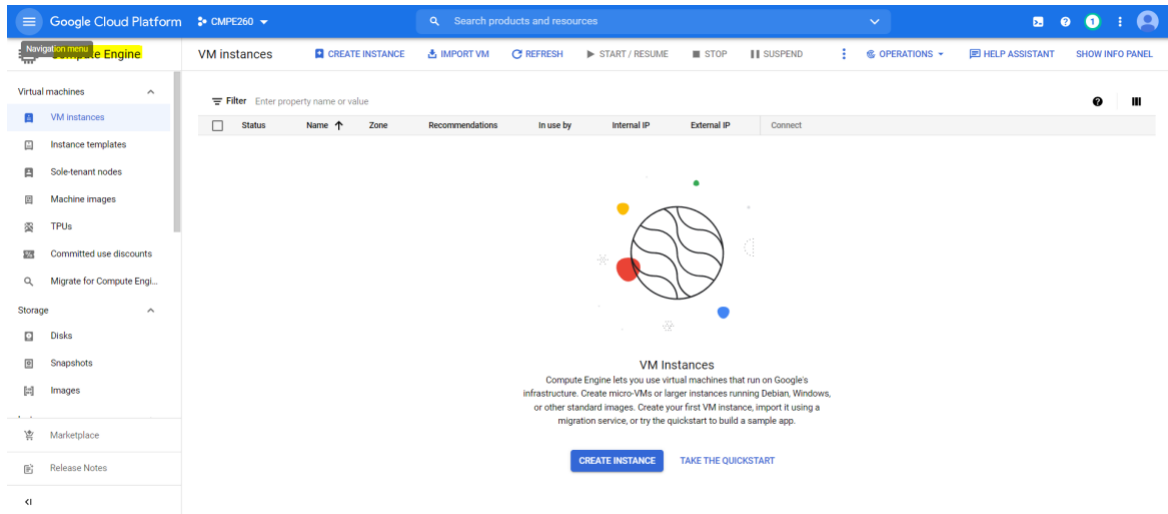
- Modify training application code for hyperparameter tuning
- Configure and launch a hyperparameter tuning job from the Vertex AI UI
- Configure and launch a hyperparameter tuning job with the Vertex SDK

Vertex AI includes many different products to support end-to-end ML workflows. This document will focus on the products highlighted below: Training/HP-Tuning and Notebooks

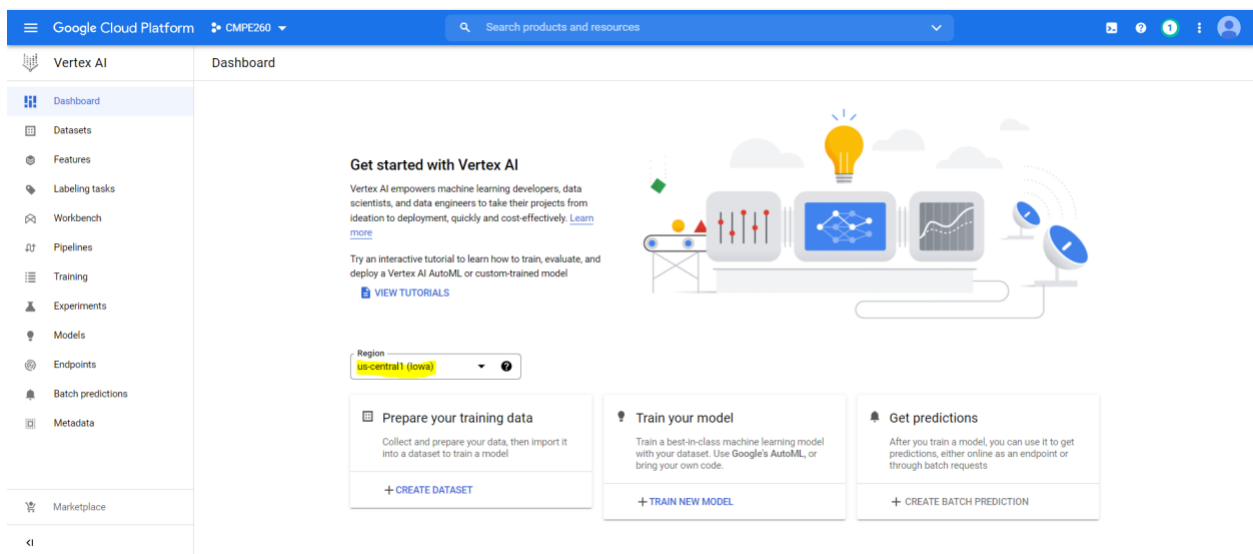


## Setup your environment

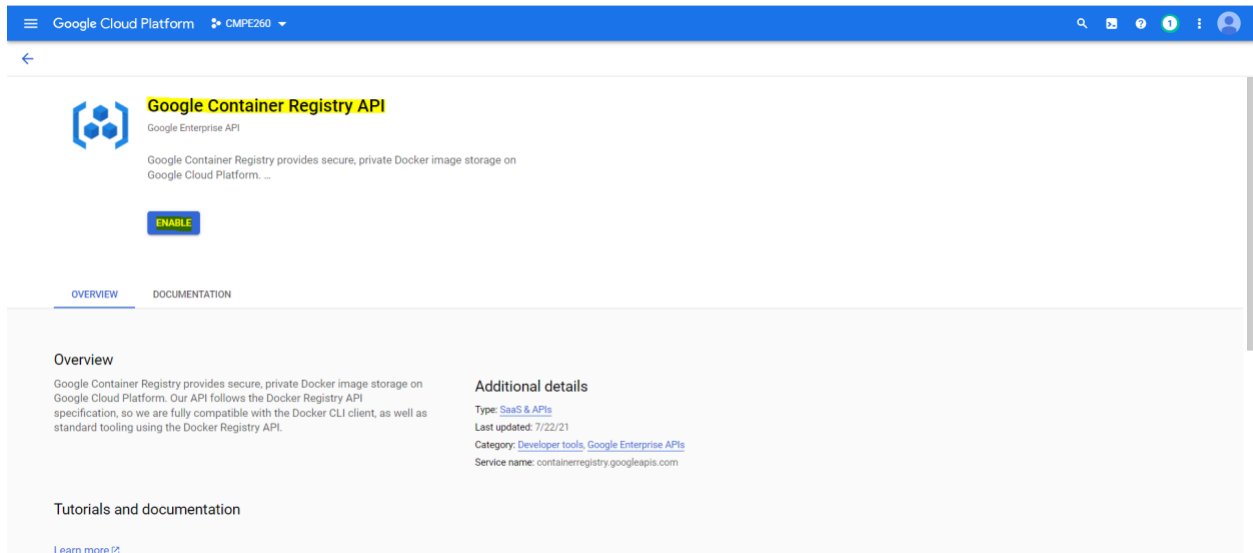
Step 1: Enable the Compute Engine API



## Step 2: Enable the Vertex AI API

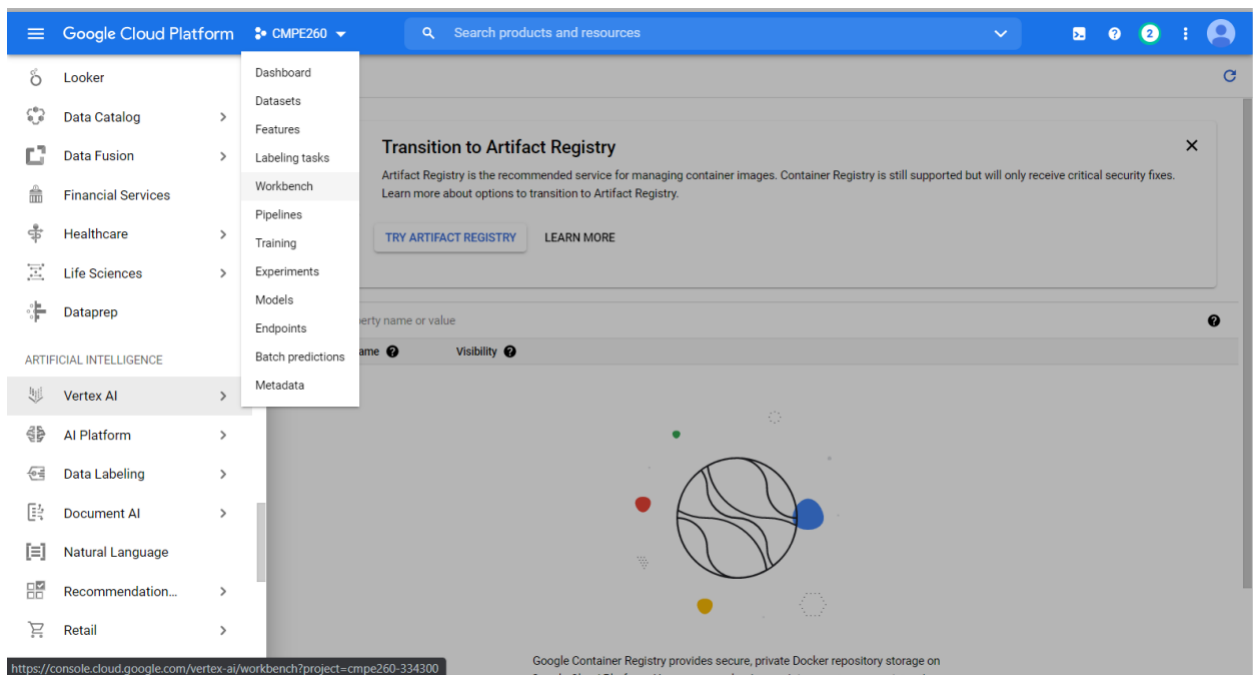


## Step 3: Enable the Container Registry API

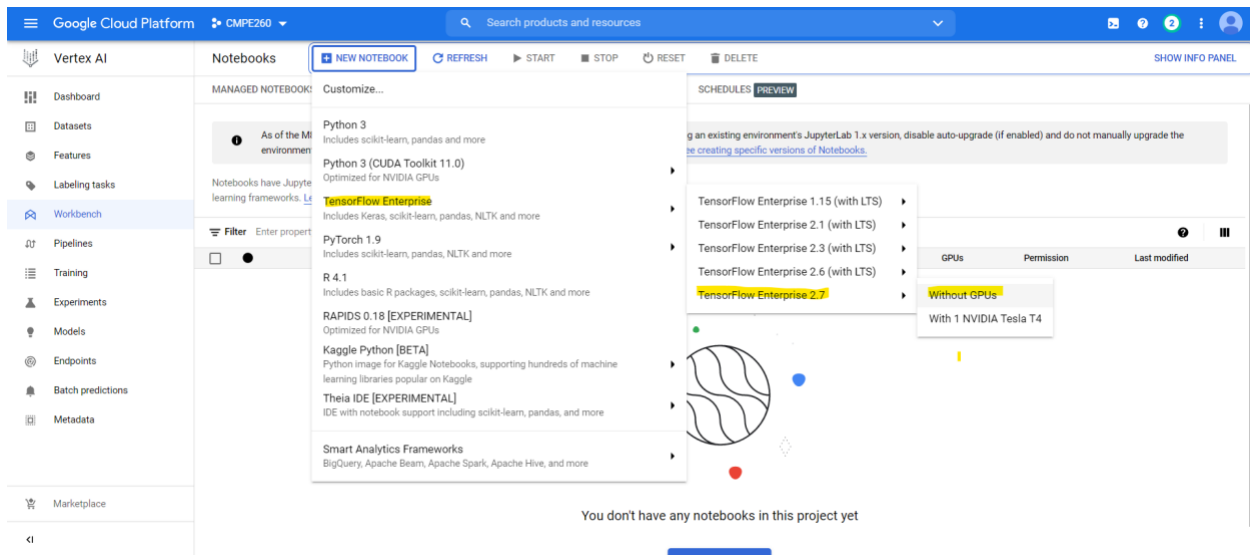


Step 4: Create a Vertex AI Workbench instance

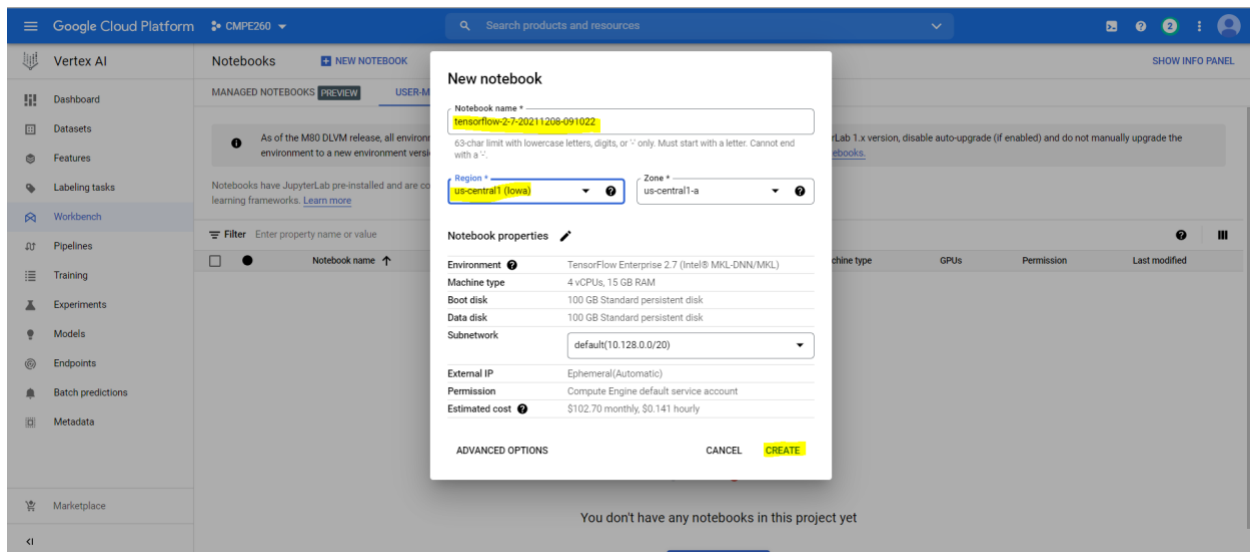
In the Left Nav, select Vertex AI → Workbench



Select **New Notebook**. Then select the **TensorFlow Enterprise 2.7** instance type **without GPUs**:



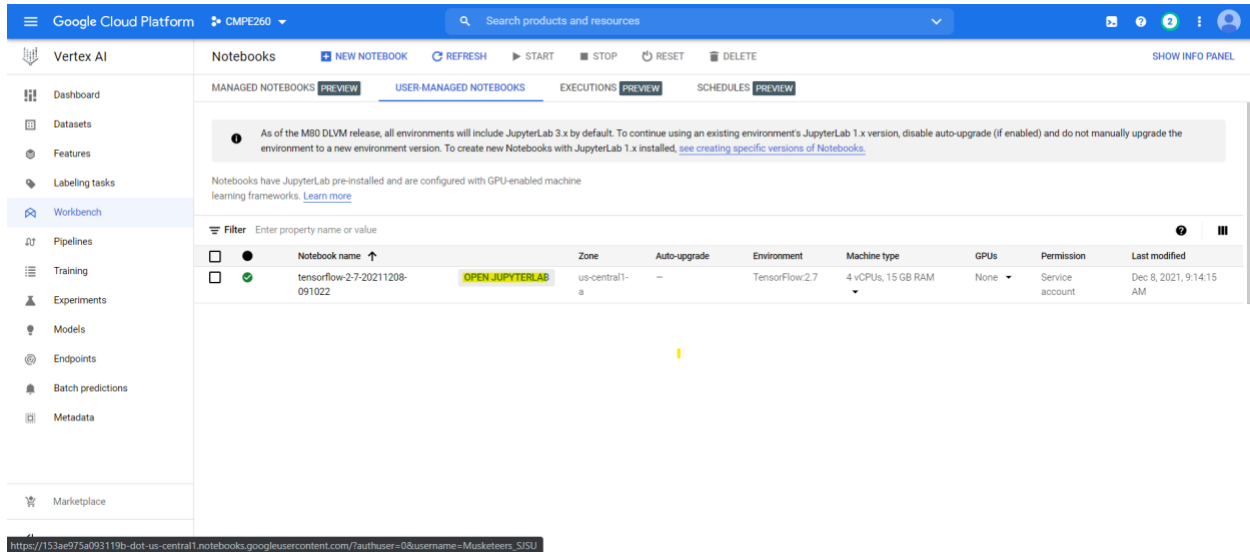
Enter the name, region and click ‘Create’



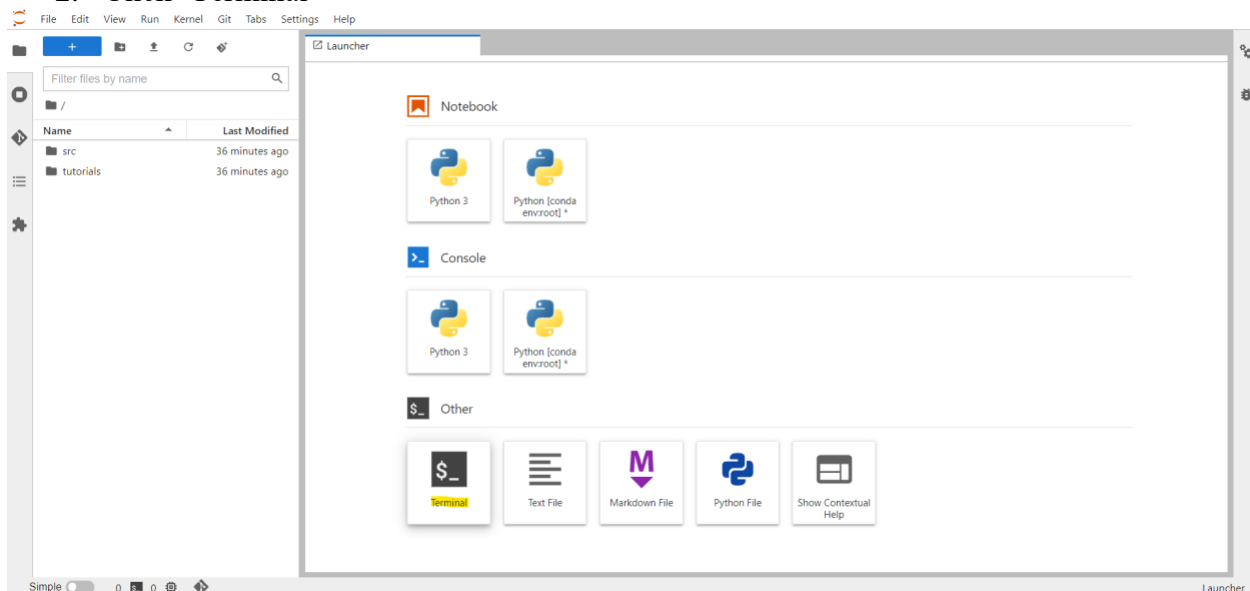
Containerize training application code

In this approach, we will submit this hyperparameter tuning job to Vertex by putting the training application code in a Docker container and pushing this container to Google Container Registry. Using this approach, you can tune hyperparameters for a model built with any framework.

### 1. Click ‘Open JupyterLab’



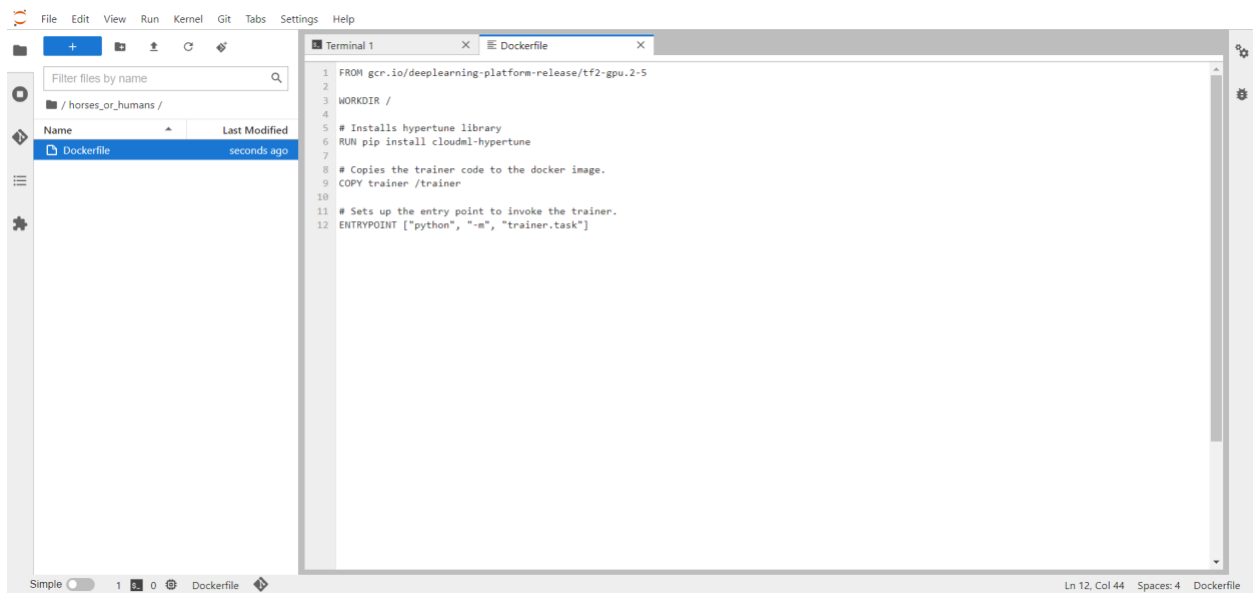

### 2. Click ‘Terminal’



Create Docker file:

Run the following commands to create a folder

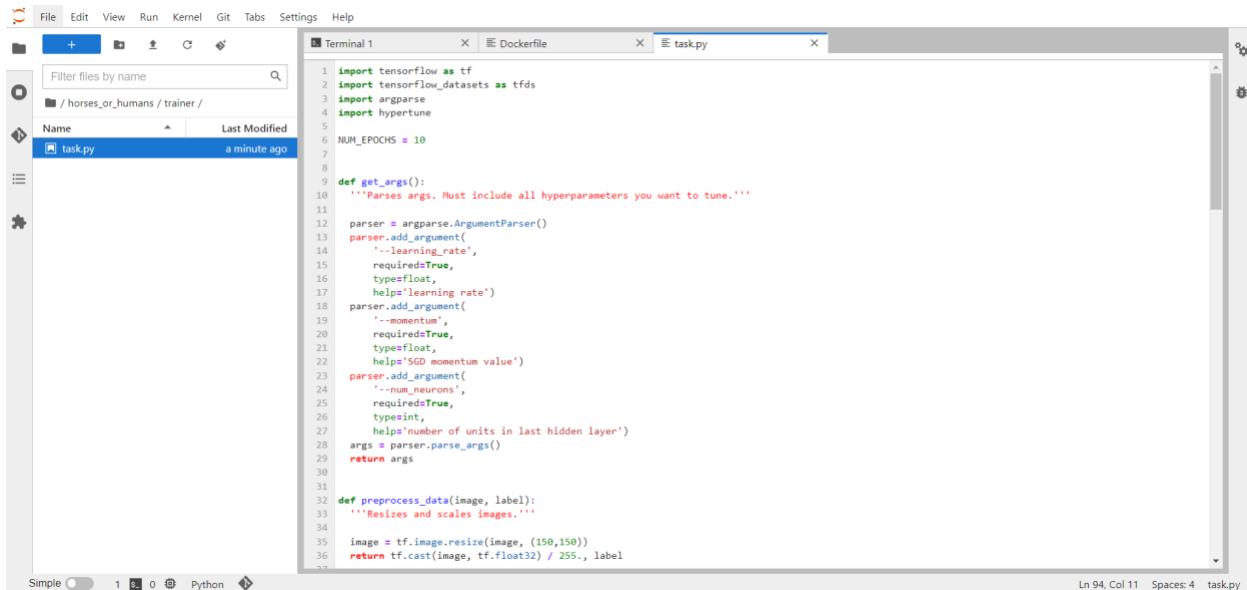
```
mkdir horses_or_humans
cd horses_or_humans
touch Dockerfile
```



*Add model training code*

Run the following commands:

```
mkdir trainer
touch trainer/task.py
```



```
1 import tensorflow as tf
2 import tensorflow_datasets as tfds
3 import argparse
4 import hypertune
5
6 NUM_EPOCHS = 10
7
8
9 def get_args():
10     """Parses args. Must include all hyperparameters you want to tune."""
11
12     parser = argparse.ArgumentParser()
13     parser.add_argument(
14         '--learning_rate',
15         required=True,
16         type=float,
17         help='learning rate')
18     parser.add_argument(
19         '--momentum',
20         required=True,
21         type=float,
22         help='SGD momentum value')
23     parser.add_argument(
24         '--num_neurons',
25         required=True,
26         type=int,
27         help='number of units in last hidden layer')
28     args = parser.parse_args()
29     return args
30
31
32 def preprocess_data(image, label):
33     """Resizes and scales images."""
34
35     image = tf.image.resize(image, (150,150))
36     return tf.cast(image, tf.float32) / 255., label
```

There are a few components that are specific to using the hyperparameter tuning service.

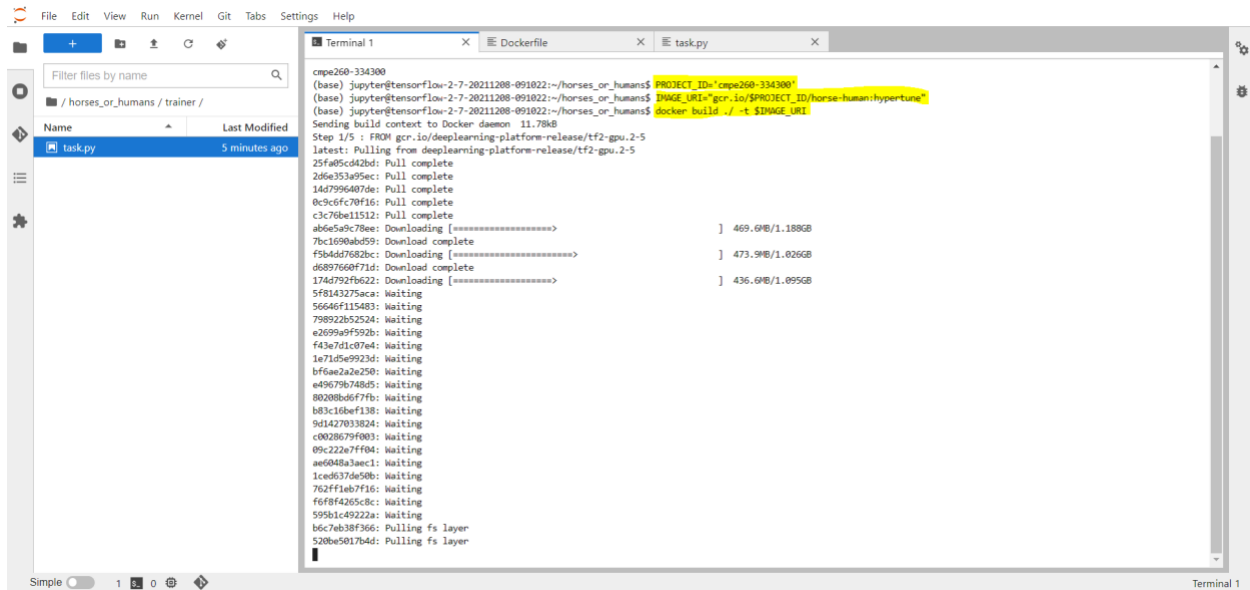
1. The script imports the hypertune library. Note that the Dockerfile from Step 1 included instructions to pip install this library.
2. The function `get_args()` defines a command-line argument for each hyperparameter you want to tune. In this example, the hyperparameters that will be tuned are the learning rate, the momentum value in the optimizer, and the number of neurons in the last hidden layer of the model, but feel free to experiment with others. The value passed in those arguments is then used to set the corresponding hyperparameter in the code.
3. At the end of the `main()` function, the hypertune library is used to define the metric you want to optimize. In TensorFlow, the `keras model.fit` method returns a `History` object. The `History.history` attribute is a record of training loss values and metrics values at successive epochs. If you pass validation data to `model.fit` the `History.history` attribute will include validation loss and metrics values as well. For example, if you trained a model for three epochs with validation data and provided accuracy as a metric, the `History.history` attribute would look similar to the following dictionary.

### *Build the container*

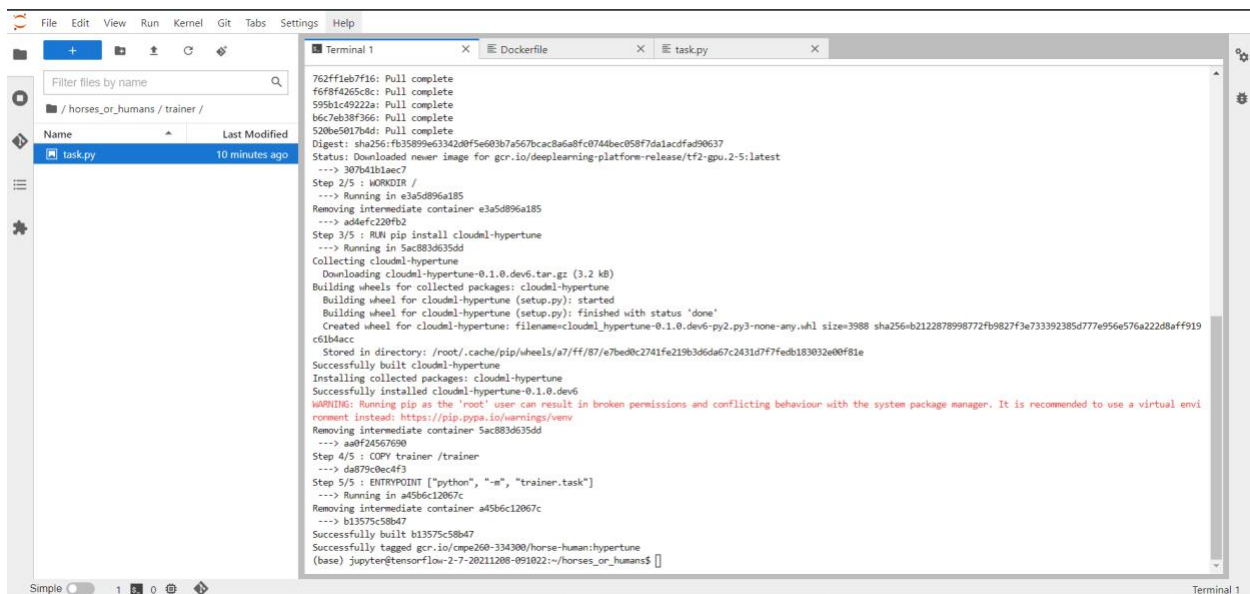
Run the following commands:

```
gcloud config list --format 'value(core.project)'
PROJECT_ID='cmpe260-334300'
```

```
IMAGE_URI="gcr.io/$PROJECT_ID/horse-human:hypertune"
docker build ./ -t $IMAGE_URI
docker push $IMAGE_URI
```

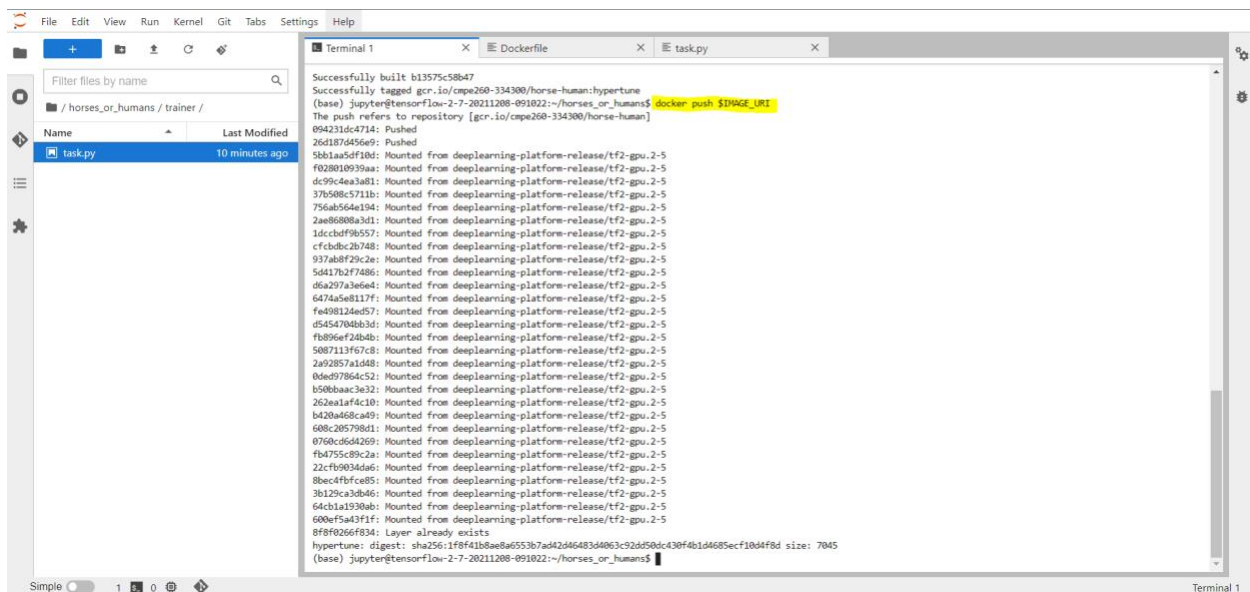


```
cmpe260-334300
(base) jupyter@tensorflow-2-7-20211208-091022:~/horses_or_humans$ PROJECT_ID="cmpe260-334300"
(base) jupyter@tensorflow-2-7-20211208-091022:~/horses_or_humans$ IMAGE_URI="gcr.io/$PROJECT_ID/horse-human:hypertune"
(base) jupyter@tensorflow-2-7-20211208-091022:~/horses_or_humans$ docker build ./ -t $IMAGE_URI
Sending build context to Docker daemon 11.76kB
Step 1/5 : FROM gcr.io/deeplearning-platform-release/tf2-gpu.2-5
latest: Pulling from deeplearning-platform-release/tf2-gpu.2-5
25fa05cd42bd: Pull complete
2d6e353a95ec: Pull complete
1a47996407de: Pull complete
8c4c6c70f16: Pull complete
c3c76be11512: Pull complete
ab6e5a9c78ee: Downloading [=====>] 469.0MB/1.188GB
7bc1690abd59: Download complete
f5b4d47682bc: Downloading [=====>] 473.9MB/1.026GB
d6897669f71d: Download complete
174d792fb622: Downloading [=====>] 436.0MB/1.095GB
5f8143275aca: Waiting
56646f115483: Waiting
798922b52524: Waiting
e3099a0f592b: Waiting
f43e7d1c074c: Waiting
1e71d5e9923d: Waiting
bf6ae2a2e250: Waiting
e49679b748d5: Waiting
80208b06f77b: Waiting
b83c16ef138: Waiting
9d1427033824: Waiting
c0028679f003: Waiting
09c222e7ff04: Waiting
ae0848a3aec1: Waiting
1cae037d650b: Waiting
762ff1eb7f16: Waiting
f6f8f4265c8c: Waiting
595b1c49222a: Waiting
b6c7eb38f366: Pulling fs layer
520be5017b4d: Pulling fs layer
```



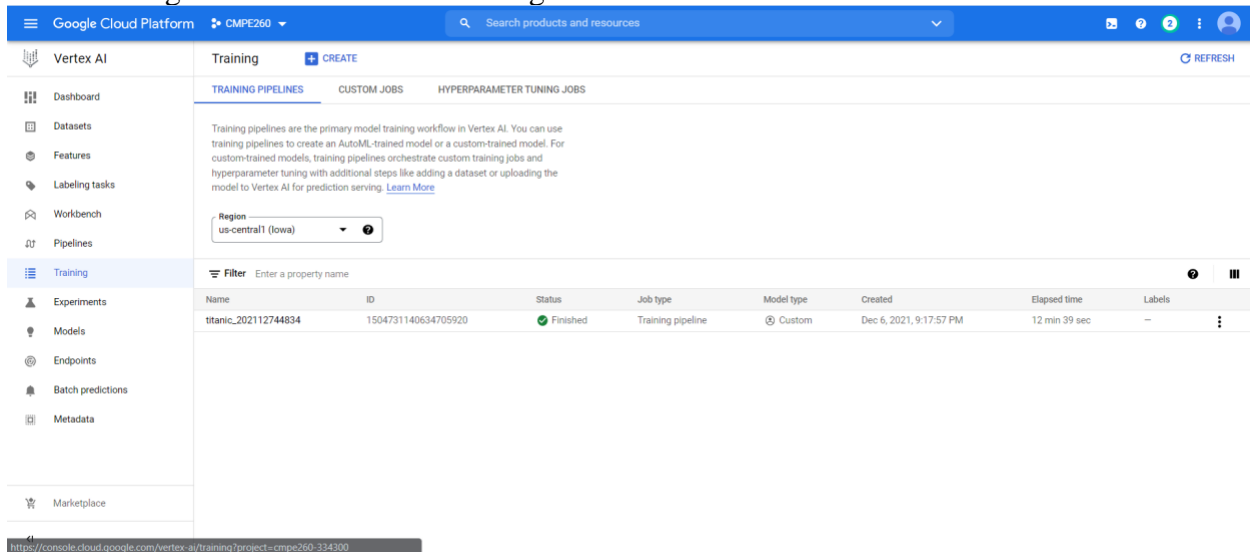
```
762ff1eb7f16: Pull complete
f6f8f4265c8c: Pull complete
595b1c49222a: Pull complete
b6c7eb38f366: Pull complete
520be5017b4d: Pull complete
Digest: sha256:fb35890e63342d0f5e683b7a567bcac8a6a8fc0744bec058f7da1acdfad90637
Status: Downloaded newer image for gcr.io/deeplearning-platform-release/tf2-gpu.2-5:latest
--> 30704b1aac7
Step 2/5 : WORKDIR /
--> Running in e3a5d896a185
Removing intermediate container e3a5d896a185
--> a4defc220fb2
Step 3/5 : RUN pip install cloudml-hypertune
--> Running in 5ac8834635dd
Collecting cloudml-hypertune
  Downloading cloudml-hypertune-0.1.0.dev6.tar.gz (3.2 kB)
Building wheels for collected packages: cloudml-hypertune
  Building wheel for cloudml-hypertune (setup.py): started
  Building wheel for cloudml-hypertune (setup.py): finished with status 'done'
  Created wheel for cloudml-hypertune: filename=cloudml_hypertune-0.1.0.dev6-py2.py3-none-any.whl size=3988 sha256=b212287898772fb9827f3e73392385d77e956e576a222d8aff919
  c61b4acc
  Stored in directory: /root/.cache/pip/wheels/a7/ff/87/e7bed0c2741fe219b3d6da67c2431d7f7fedb183032e0f81e
Successfully built cloudml-hypertune
Installing collected packages: cloudml-hypertune
Successfully installed cloudml-hypertune-0.1.0.dev6
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
Removing intermediate container 5ac8834635dd
--> a4f24567698
Step 4/5 : COPY trainer/trainer
--> da879c0ec4f3
Step 5/5 : ENTRYPOINT ["python", "-m", "trainer.task"]
--> Running in a456c12867c
Removing intermediate container a456c12867c
--> b13575c58b47
Successfully built b13575c58b47
Successfully tagged gcr.io/cmpe260-334300/horse-human:hypertune
(base) jupyter@tensorflow-2-7-20211208-091022:~/horses_or_humans$ []
```





# Run a hyperparameter tuning job on Vertex AI

## 1. Navigate to Vertex AI → Training



## Step 1: Configure training job

Click **Create** to enter the parameters for your hyperparameter tuning job.

- Under **Dataset**, select **No managed dataset**
- Then select **Custom training (advanced)** as your training method and click **Continue**.
- Click **Continue**

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

1 Training method

2 Model details

3 Training container

4 Hyperparameters (optional)

5 Compute and pricing

6 Prediction container (optional)

START TRAINING CANCEL

Dataset \*

No managed dataset

Annotation set

Objective

Custom

Please refer to the pricing guide for more details (and available deployment options) for each method.

AutoML options are only available when you train with a managed dataset.

AutoML

Train high-quality models with minimal effort and machine learning expertise. Just specify how long you want to train. [Learn more](#)

AutoML Edge

Train a model that can be exported for on-prem/on-device use. Typically has lower accuracy. [Learn more](#)

Custom training (advanced)

Run your TensorFlow, scikit-learn, and XGBoost training applications in the cloud. Train with one of Google Cloud's pre-built containers or use your own. [Learn more](#)

CONTINUE

- Enter **horses-humans-hyptertune** (or whatever you'd like to call your model) for **Model name**
- Click **Continue**

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

1 Training method

2 Model details

3 Training container

4 Hyperparameters (optional)

5 Compute and pricing

6 Prediction container (optional)

START TRAINING CANCEL

Model name \*

horses-humans-hyptertune

Encryption

Use a customer-managed encryption key (CMEX)

Service account

Select a service account to use with your model.

Service account BROWSE

Network

The full name of the Compute Engine network to which the job should be peered.

Peered VPC network

Training Debugging

The interactive terminal enables interactive debugging and profiling.

Enable training debugging

SHOW LESS

CONTINUE

- Select 'Custom Container'
- Enter **Container Image: URI to the Image**

gcr.io/cmpe260-334300/horse-human@sha256:1f8f41b8ae8a6553b7ad42d46483d4063c92dd50dc430f4b1d4685ecf10d4f8d

The screenshot shows the 'Train new model' wizard in the Google Cloud Platform console. The left sidebar lists various services, with 'Training' selected. The main panel shows the 'Training container' step, which is highlighted. The 'Hyperparameters (optional)' step is also visible. The 'Custom container' option is selected under 'Pre-built container'. The 'Container image' field contains the SHA256 hash of the Docker image. The 'Model output directory' field is empty. The 'Arguments' field contains the command line arguments: `-flag_a=xxxxx`, `-flag2`, and `flag3`. The 'CONTINUE' button is visible at the bottom.

## Step 2: Configure hyperparameter tuning job

- Select 'Enable Hyperparameter tuning'
- Enter the Parameter details for `learning_rate`

The screenshot shows the 'Train new model' wizard in the Google Cloud Platform console, specifically the 'Hyperparameters (optional)' step. The 'Enable hyperparameter tuning' checkbox is checked. The 'Hyperparameter tuning variables' section is visible, showing a 'New Hyperparameter' form. The 'Parameter name' field contains `learning_rate`. The 'Type' is set to 'Double'. The 'Min' value is 0.01 and the 'Max' value is 1. The 'Scaling' is set to 'Log'. The 'ADD NEW PARAMETER' button is visible at the bottom.

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

Training method

Model details

Training container

Hyperparameters (optional)

Compute and pricing

START TRAINING CANCEL

Hyperparameter tuning optimizes your model through multiple trials in one training job, but will increase the cost of this job. After training finishes, the best-performing model will be saved to your Model List. [Learn more](#)

☒ Enable hyperparameter tuning

### Hyperparameter tuning variables

1 Ensure that your hyperparameter variables are named and typed correctly. [VIEW DOCS](#)

learning\_rate (Double), 0.01 - 1

New Hyperparameter

Parameter name \*  
momentum

Type \*  
Double

Min \*  
0

Max \*  
1

Scaling \*  
Linear

CANCEL DONE

ADD NEW PARAMETER

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

Training method

Model details

Training container

Hyperparameters (optional)

Compute and pricing

START TRAINING CANCEL

☒ Enable hyperparameter tuning

### Hyperparameter tuning variables

1 Ensure that your hyperparameter variables are named and typed correctly. [VIEW DOCS](#)

learning\_rate (Double), 0.01 - 1

momentum (Double), 0 - 1

New Hyperparameter

Parameter name \*  
num\_neurons

Type \*  
Discrete

Values \*  
64,128,512

Enter all the values you want to tune, separated with comma.

Scaling \*  
No scaling

CANCEL DONE

ADD NEW PARAMETER

Metric to optimize \*

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

Training method

Model details

Training container

Hyperparameters (optional)

Compute and pricing

START TRAINING CANCEL

Hyperparameter tuning variables

Ensure that your hyperparameter variables are named and typed correctly.

VIEW DOCS

learning\_rate (Double): 0.01 - 1

momentum (Double): 0 - 1

num\_neurons (Discrete): 64128, 512

ADD NEW PARAMETER

Metric to optimize \*

accuracy

Goal \*

Maximize

Maximum number of trials \*

15

How many training trials should be attempted to optimize the specified hyperparameters. Increasing the number of trials generally yields better results but also increases cost.

Learn more

Maximum number of parallel trials \*

4

The number of training trials to run concurrently. More parallel trials shortens training time but reduces the effectiveness of the tuning.

Algorithm \*

Default

Search algorithms for hyperparameter tuning.

CONTINUE

## Step 3: Configure compute

Upgrade your account to avoid a break in service

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Train new model

Training method

Model details

Training container

Hyperparameters (optional)

Compute and pricing

START TRAINING CANCEL

Model training pricing is based on the length of time spent training, machine types, and any accelerators used. [Learn more](#)

Region

us-central1 (Iowa)

Compute settings

Select the type of virtual machine to use for your worker pool. You can add up to 4 worker pools. To learn about compute costs and how to map your ML framework's roles to specific worker pools, consult the [documentation](#)

Worker pool 0

Machine type \*

n1-standard-4, 4 vCPUs, 15 GiB memory

Accelerator type

NVidia, TESLA, T4

Accelerators can speed up model training that involves intensive compute tasks. [Learn more](#)

Accelerator count

1

Worker count

1

Disk type

SSD

Disk size (GB)

100

ADD MORE WORKER POOLS (OPTIONAL)

## Click 'Start Tuning'

Upgrade your account to avoid a break in service (\$159.65 credit and 6 days left in your trial).

Google Cloud Platform

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

### Train new model

- Training method
- Model details
- Training container
- Hyperparameters (optional)
- Compute and pricing**

**START TRAINING** CANCEL

Model training pricing is based on the length of time spent training, machine types, and any accelerators used. [Learn more](#)

Region: us-central1 (Iowa)

#### Compute settings

Select the type of virtual machine to use for your worker pool. You can add up to 4 worker pools. To learn about compute costs and how to map your ML framework's roles to specific worker pools, consult the [documentation](#).

#### Worker pool 0

Machine type: n1-standard-4, 4 vCPUs, 15 GiB memory

Accelerator type: NVIDIA\_TESLA\_T4

Accelerators can speed up model training that involves intensive compute tasks. [Learn more](#)

Accelerator count: 1

Worker count: 1

Disk type: SSD

Disk size (GB): 100

[ADD MORE WORKER POOLS \(OPTIONAL\)](#)

## HyperParameter Tuning job is in progress:

Upgrade your account to avoid a break in service (\$159.65 credit and 6 days left in your trial).

Google Cloud Platform

Training

TRAINING PIPELINES CUSTOM JOBS **HYPERPARAMETER TUNING JOBS**

Hyperparameter tuning searches for the best combination of hyperparameter values by optimizing metric values across a series of trials. Hyperparameter tuning is only used by custom-trained models and not AutoML models. [Learn More](#)

Region: us-central1 (Iowa)

Filter: Enter a property name

| Name  | ID                  | Status         | Job type                  | Model type | Created                  | Elapsed time | Labels |
|---|---------------------|----------------|---------------------------|------------|--------------------------|--------------|--------|
| <b>horses-humans-hyptertune-hyperparameter-tuning-job</b> | 2213537908550270976 | <b>Running</b> | Hyperparameter tuning job | —          | Dec 8, 2021, 10:24:04 AM | 19 sec       | —      |

# Cleanup

## Stop the notebook instance

Google Cloud Platform

CMPE260

Search products and resources

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Notebooks

NEW NOTEBOOK

REFRESH

START

STOP

RESET

DELETE

SHOW INFO PANEL

MANAGED NOTEBOOKS

USER-MANAGED NOTEBOOKS

EXECUTIONS

SCHEDULES

As of the M80 DLM release, all environments will include JupyterLab 3.x by default. To continue using an existing environment's JupyterLab 1.x version, disable auto-upgrade (if enabled) and do not manually upgrade the environment to a new environment version. To create new Notebooks with JupyterLab 1.x installed, [see creating specific versions of Notebooks](#).

Notebooks have JupyterLab pre-installed and are configured with GPU-enabled machine learning frameworks. [Learn more](#)

Filter

Enter property name or value

|                                     | Notebook name                  | Zone            | Auto-upgrade  | Environment    | Machine type       | GPUs | Permission      | Last modified           |
|-------------------------------------|--------------------------------|-----------------|---------------|----------------|--------------------|------|-----------------|-------------------------|
| <input checked="" type="checkbox"/> | tensorflow-2.7-20211208-091109 | OPEN JUPYTERLAB | us-central1-a | TensorFlow.2.7 | 4 vCPUs, 15 GB RAM | None | Service account | Dec 8, 2021, 9:14:15 AM |

## Delete the Storage Bucket

Google Cloud Platform

CMPE260

Search products and resources

Cloud Storage

Browser

Monitoring

Settings

Release Notes

Browser

CREATE BUCKET

DELETE

REFRESH

SHOW INFO PANEL

Filter

Filter buckets

| Name   | Created                  | Location type | Location           | Default storage class | Last modified            | Public access          | Access control |
|--|--------------------------|---------------|--------------------|-----------------------|--------------------------|------------------------|----------------|
| <input checked="" type="checkbox"/> artifacts-cmpe260-3343303.appspot.c... | Dec 8, 2021, 10:07:17 AM | Multi-region  | us (multiple re... | Standard              | Dec 8, 2021, 10:07:17 AM | Subject to object ACLs | Fine-grained   |
| <input type="checkbox"/> cloud-ai-platform-65473780-0280-46f...            | Dec 6, 2021, 4:17:06 PM  | Region        | us-central1 (lo... | Regional              | Dec 6, 2021, 4:17:06 PM  | Subject to object ACLs | Fine-grained   |